

Notes of DP - Bayesian Inference

1 Bayesian Inference Based on Beta-Bernoulli Distribution

In the Bayesian inference, first there is a prior distribution $\pi(\alpha)$ to present our belief about parameter α . Then, we get some observed data x sizing n , and produce a posterior distribution $Pr(\alpha|x)$. The Bayesian inference is based on the Bayes' rule to calculate the posterior distribution:

$$Pr(\alpha|x) = \frac{Pr(x|\alpha)\pi(\alpha)}{Pr(x)}$$

It is denoted as $Bl(x, \pi(\alpha))$ taking an observed data set $x \in \mathcal{X}^n$ and a prior distribution $\pi(\alpha)$ as input, outputting a posterior distribution *posterior*. For conciseness, when prior is given, we use $Bl(x)$. n is the size of the observed data size.

In our inference algorithm, we take a Dirichlet distribution as *prior*, $DL(\alpha_1, \alpha_2, \dots, \alpha_m)$. Columns of the data set is corresponding to the number of parameters (m) in Dirichlet distribution. The posterior distribution is still a Dirichlet distribution with parameters $(\alpha_1 + n_1, \alpha_2 + n_2, \dots, \alpha_m + n_m)$, where n_i is the number of 1 at the column i in observed data set.

In the case that $m = 2$, it is reduced to a Beta distribution $\text{beta}(\alpha, \beta)$. We can have only one column take value from $\{0, 1\}$. The posterior is computed as $(\alpha + n_1, \beta + n_0)$, where n_1 is the number of 1 in the data set and n_0 is the number of 0 in the data set.

2 Algorithm Setting up

For now, we already have a prior distribution *prior*, an observed data set x .

2.1 Exponential Mechanism with Global Sensitivity

In exponential mechanism, candidate set R can be obtained by enumerating $y \in \mathcal{X}^n$, i.e.

$$R = \{Bl(y) \mid y \in \mathcal{X}^n\}.$$

Hellinger distance H is used here to score these candidates. The utility function:

$$u(x, r) = -H(Bl(x), r); r \in R. \quad (1)$$

Exponential mechanism with global sensitivity selects and outputs a candidate $r \in R$ with probability proportional to $\exp(\frac{\epsilon u(x, r)}{2\Delta_g u})$:

$$P[r] = \frac{\exp(\frac{\epsilon u(x, r)}{2\Delta_g u})}{\sum_{r' \in R} \exp(\frac{\epsilon u(x, r')}{2\Delta_g u})},$$

where global sensitivity is calculated by:

$$\Delta_g u = \max_{\{|x', y'| \leq 1; x', y' \in \mathcal{X}^n\}} \max_{\{r \in R\}} |H(Bl(x'), r) - H(Bl(y'), r)|$$

The basic exponential mechanism is ϵ -differential privacy[1].

2.2 Exponential Mechanism with Local Sensitivity

Exponential mechanism with local sensitivity share the same candidate set and utility function as it with global sensitivity. This outputs a candidate $r \in R$ with probability proportional to $\exp(\frac{\epsilon u(x,r)}{2\Delta_l u})$:

$$P[r] = \frac{\exp(\frac{\epsilon u(x,r)}{2\Delta_l u})}{\sum_{r' \in R} \exp(\frac{\epsilon u(x,r')}{2\Delta_l u})},$$

where local sensitivity is calculated by:

$$\Delta_l u(x) = \max_{\{x, y' | |x - y'| \leq 1; y' \in \mathcal{X}^n\}} \max_{\{r \in R\}} |H(\text{BI}(x), r) - H(\text{BI}(y'), r)|$$

The exponential mechanism with local sensitivity is non differential privacy[1].

2.3 Exponential Mechanism with Smooth Sensitivity

2.3.1 Algorithm Setting up

The candidate set and utility function are still the same as before, differ only in the sensitivity. It will output a candidate $r \in R$ with probability proportional to $\exp(\frac{\epsilon u(x,r)}{2S(x)})$:

$$P[r] = \frac{\exp(\frac{\epsilon u(x,r)}{2S(x)})}{\sum_{r' \in R} \exp(\frac{\epsilon u(x,r')}{2S(x)})},$$

where the sensitivity in mechanism is smooth sensitivity $S(x)$ [2], calculated by:

$$S_\beta(x) = \max(\Delta_l u(x), \max_{y \neq x; y \in D^n} (\Delta_l u(y) \cdot e^{-\beta d(x,y)})),$$

where $\beta = \beta(\epsilon, \delta)$. In our private Bayesian inference mechanism, we set the β as $\ln(1 - \frac{\epsilon}{2 \ln(\frac{\delta}{2(n+1)})})$.

2.3.2 Sliding Property of Exponential Mechanism

Lemma 2.1. *for any exponential mechanism $\mathcal{M}_E(x, u, \mathcal{R})$, $\lambda = f(\epsilon, \delta)$, ϵ and $|\delta| < 1$, the sliding property holds:*

$$\Pr_{z \sim \mathcal{M}_E(x, u, \mathcal{R})} [u(r, x) = \hat{s}] \leq e^{\frac{\epsilon}{2}} \Pr_{z \sim \mathcal{M}_E(x, u, \mathcal{R})} [u(r, x) = (\Delta + \hat{s})] + \frac{\delta}{2},$$

Proof. We denote the normalizer of the probability mass in $\mathcal{M}_E(x, u, \mathcal{R})$: $\sum_{r' \in \mathcal{R}} \exp(\frac{\epsilon u(r', x)}{2S(x)})$ as NL_x :

$$\begin{aligned} LHS &= \Pr_{z \sim \mathcal{M}_E(x, u, \mathcal{R})} [u(r, x) = \hat{s}] = \frac{\exp(\frac{\epsilon \hat{s}}{2S(x)})}{NL_x} \\ &= \frac{\exp(\frac{\epsilon(\hat{s} + \Delta - \Delta)}{2S(x)})}{NL_x} \\ &= \frac{\exp(\frac{\epsilon(\hat{s} + \Delta)}{2S(x)} + \frac{-\epsilon \Delta}{2S(x)})}{NL_x} \\ &= \frac{\exp(\frac{\epsilon(\hat{s} + \Delta)}{2S(x)})}{NL_x} \cdot e^{\frac{-\epsilon \Delta}{2S(x)}}. \end{aligned}$$

By bounding the $\Delta \geq -S(x)$, we can get:

$$\begin{aligned} \frac{\exp(\frac{\epsilon(\hat{s}+\Delta)}{2S(x)})}{NL_x} \cdot e^{\frac{-\epsilon\Delta}{2S(x)}} &\leq \frac{\exp(\frac{\epsilon(\hat{s}+\Delta)}{2S(x)})}{NL_x} \cdot e^{\frac{\epsilon}{2}} \\ &= e^{\frac{\epsilon}{2}} \Pr_{z \sim \mathcal{M}_E(x, u, \mathcal{R})} [u(r, x) = (\Delta + \hat{s})] \leq RHS \end{aligned}$$

□

2.3.3 Dilation Property of Exponential Mechanism

Lemma 2.2. *for any exponential mechanism $\mathcal{M}_E(x, u, \mathcal{R})$, $\lambda < |\beta|$, ϵ , $|\delta| < 1$ and $\beta \leq \ln(1 - \frac{\epsilon}{2 \ln(\frac{\delta}{2(n+1)})})$, the dilation property holds:*

$$\Pr_{r \sim \mathcal{M}_E(x, u, \mathcal{R})} [u(r) = z] \leq e^{\frac{\epsilon}{2}} \Pr_{r \sim \mathcal{M}_E(x, u, \mathcal{R})} [u(r) = e^\lambda z] + \frac{\delta}{2},$$

where the sensitivity in mechanism is still smooth sensitivity as above.

Proof. The sensitivity is always greater than 0, and we are using $-H(\text{Bl}(x), r)$ for utility function, i.e., $u(r) \leq 0$, we need to consider two cases that $\lambda < 0$, and $\lambda > 0$:

We set the $h(z) = \Pr[u(\mathcal{M}_E(x, u, \mathcal{R})) = z] = \frac{\exp(\frac{\epsilon z}{2S(x)})}{NL_x}$.

We first consider $\lambda < 0$. In this case, $1 < e^\lambda$, so the ratio $\frac{h(z)}{h(e^\lambda z)} = \frac{\exp(\frac{\epsilon z}{2S(x)})}{\exp(\frac{\epsilon(e^\lambda z)}{2S(x)})}$ is at most $\frac{\epsilon}{2}$.

Next, we proof the dilation property for $\lambda > 0$, The ratio of $\frac{h(z)}{h(e^\lambda z)}$ is $\exp(\frac{\epsilon}{2} \cdot \frac{u(\mathcal{M}_E(x, u, \mathcal{R}))(1-e^\lambda)}{S(x)})$. Consider the event $G = \{\mathcal{M}_E(x, u, \mathcal{R}) : u(\mathcal{M}_E(x, u, \mathcal{R})) \leq \frac{S(x)}{(1-e^\lambda)}\}$. Under this event, the log-ratio above is at most $\frac{\epsilon}{2}$. The probability of G under density $h(z)$ is $1 - \frac{\delta}{2}$. Thus, the probability of a given event z is at most $\Pr[z \cap G] + \Pr[\bar{G}] \leq e^{\frac{\epsilon}{2}} \Pr[e^\lambda z \cap G] + \frac{\delta}{2} \leq e^{\frac{\epsilon}{2}} \Pr[e^\lambda z] + \frac{\delta}{2}$.

Detail proof:

- $\lambda < 0$

The left hand side will always be smaller than 0 and the right hand side greater than 0. This will always holds, i.e.

- $\lambda > 0$

Because $\hat{s} = u(r)$ where $r \sim \mathcal{M}_E(x, u, \mathcal{R})$, we can substitute \hat{s} with $u(\mathcal{M}_E(x, u, \mathcal{R}))$. Then, what we need to proof under the case $\lambda > 0$ is:

$$u(\mathcal{M}_E(x, u, \mathcal{R})) \leq \frac{S(x)}{(1-e^\lambda)}$$

By applying the accuracy property of exponential mechanism, we bound the probability that the equation holds with probability:

$$\Pr[u(\mathcal{M}_E(x, u, \mathcal{R})) \leq \frac{S(x)}{(1-e^\lambda)}] \leq \frac{|\mathcal{R}| \exp(\frac{\epsilon S(x)}{(1-e^\lambda)}/2S(x))}{|\mathcal{R}_{OPT}| \exp(\epsilon OPT_{u(x)}/2S(x))}$$

In our Bayesian Inference mechanism, the size of the candidate set \mathcal{R} is equal to the size of observed data set plus 1, i.e., $n + 1$, and $OPT_{u(x)} = 0$, then we have:

$$\begin{aligned} \Pr[u(\mathcal{M}_E(x, u, \mathcal{R})) \leq \frac{S(x)}{(1-e^\lambda)}] &= (n+1) \exp(\frac{\epsilon S(x)}{(1-e^\lambda)}/2S(x)) \\ &= (n+1) \exp(\frac{\epsilon}{2(1-e^\lambda)}) \end{aligned}$$

When we set $\lambda \leq \ln(1 - \frac{\epsilon}{2\ln(\frac{\delta}{2(n+1)})})$, it is easily to derive that $Pr[u(\mathcal{M}_E(x, u, \mathcal{R})) \leq \frac{S(x)}{(1-e^\lambda)}] \leq \frac{\delta}{2}$.

□

3 Experimental Evaluations

3.1 Computation Efficiency

The formula for computing the local sensitivity is presented in Sec. 2.2: $\max_{\{|x, y'| \leq 1; y' \in \mathcal{X}^n\}} \max_{\{r \in R\}} \{H(\text{Bl}(x), r) - H(\text{Bl}(y'), r)\}$ can be reduced to $\max_{\{|x, y'| \leq 1; y' \in \mathcal{X}^n\}} H(\text{Bl}(x), \text{Bl}(y'))$ by applying the distance triangle property. i.e., the maximum value over $\max_{r \in R}$ always happen when $r = \text{Bl}(x)$ itself, where $\Delta_l u(x) = \max_{\{|x, y'| \leq 1; y' \in \mathcal{X}^n\}} \{H(\text{Bl}(x), \text{Bl}(x)) - H(\text{Bl}(y'), \text{Bl}(x))\} = \max_{\{|x, y'| \leq 1; y' \in \mathcal{X}^n\}} \{H(\text{Bl}(y'), \text{Bl}(x))\}$. We also have some experiments for validating our proposal as in Fig. 3.1, where we calculate the $\max_{\{|x, y'| \leq 1; y' \in \mathcal{X}^n\}}$ value for every candidate $r \in R$. It is shown that maximum value taken when $r = \text{Bl}(x)$.




Figure 1: Experimental Results for Finding the Local Sensitivity Efficiently

3.2 Accuracy Study with Hellinger Distance

We study the accuracy property of the exponential mechanism with three kind of sensitivity, as well as the Laplace mechanism for comparison. As in Fig. 2, the two figures shown the accuracy of four algorithms under same configuration. The accuracy is measured by Hellinger distance between output r^* and the true inference result $\text{Bl}(x)$, $H(r^*, \text{Bl}(x))$.

References

- [1] Cynthia Dwork, Aaron Roth, et al. “The algorithmic foundations of differential privacy”. In: *Foundations and Trends® in Theoretical Computer Science* 9.3–4 (2014), pp. 211–407.

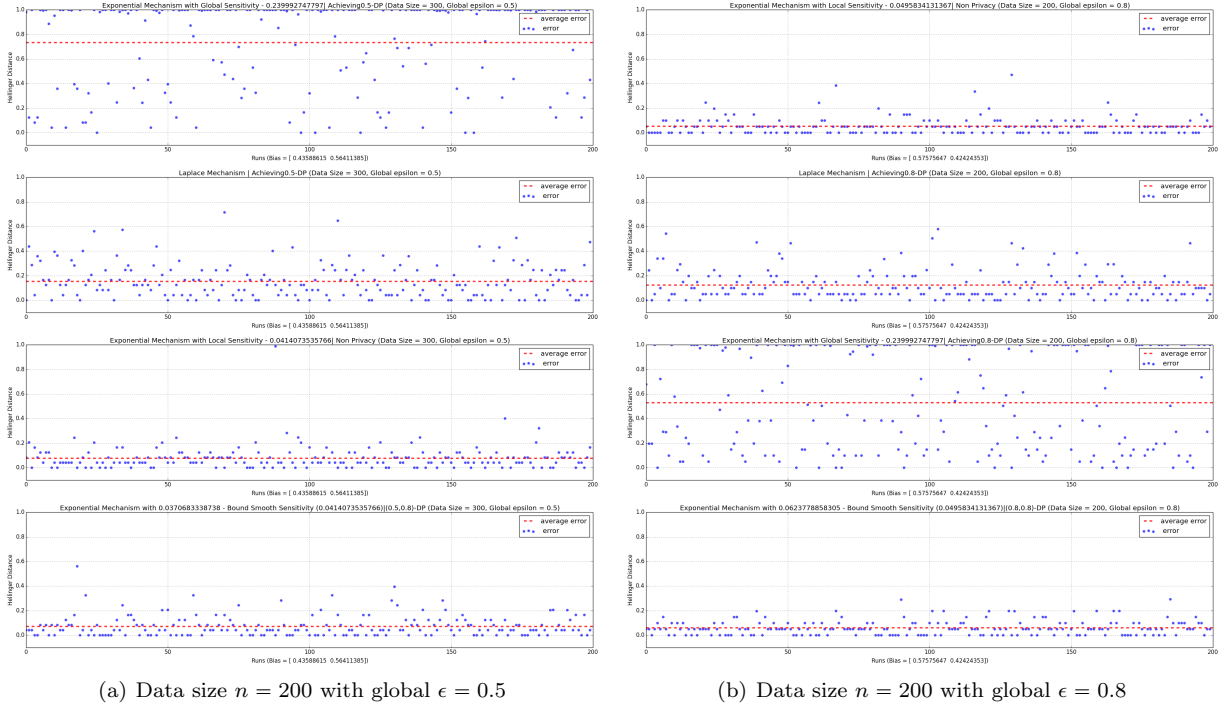


Figure 2: The experimental results of accuracy of four algorithms

- [2] Kobbi Nissim, Sofya Raskhodnikova, and Adam Smith. “Smooth sensitivity and sampling in private data analysis”. In: *Proceedings of the thirty-ninth annual ACM symposium on Theory of computing*. ACM, 2007, pp. 75–84.